

COMP6714 2024T3 Project: Ranked Retrieval with Spelling Correction

Please read this spec carefully and completely before you start programming.

In this project, you are going to implement (using Python3 in CSE linux machines) a simple search engine that ranks the output documents based on the proximity of the matching terms. It also supports spelling correction of query terms with maximum editing distance of 2 per search term (assuming Insert, Delete and Replace operations and no transpose). A search query in this project is a list of space-separated search terms and each search term may contain any numeric digits or uppercase/lowercase letters, and will not contain any punctuations. You will need to implement an indexer (to index the files) and a search program (to search the files based on the index generated by your indexer). As a core requirement for this project, you must implement your solution using an inverted index with positional information (for example, the positional index described in Lecture Week 1). You may also implement any additional indexes as appropriate, if you wish.

Given a search query, each matching document from the search result must contain terms that match all the search terms following the below term matching rules:

- Search is case insensitive.
- Full stops for abbreviations are ignored. e.g., U.S., US are the same.
- Singular/Plural/Possessive is ignored. e.g., cat, cats, cat's, cats' are all the same.
- Tense is ignored. e.g., breaches, breach, breached, breaching are all the same.
- Numeric tokens such as years, integers should be indexed accordingly and searchable.
- Commas in numeric tokens are ignored, e.g., 1,000,000 and 1000000 are the same.
- Numbers with decimal places are ignored from the index, as a decimal number is not a valid search term (since '.' is not allowed).
- Except the above, all other punctuation should be treated as token dividers.

As a requirement of this project, the matching documents in a search result are ranked according to the distances between the matching terms in these documents, such that matching terms closer to each other will be ranked higher than those further apart. Further details are described below in the Ranking section.

You are provided with approximately 1000 small documents (named with their document IDs) available in `~cs6714/reuters/data`. You can find these files by logging into CSE machines and going to folder `~cs6714/reuters/data`. Your submitted project will be tested against a similar collection of up to 1000 documents (i.e., we may replace some of these documents to avoid any hard-coded solutions).

Your submission must include 2 main programs: `index.py` and `search.py` as described below. You may submit additional Python files in addition to these 2 files. It is your responsibility to submit any other required Python files for the 2 main programs to work properly.

This project will be marked based on auto marking, and will then be checked manually for other requirements described in this specification (for example, if a positional index is implemented). To ensure your project satisfies the input and output formatting requirements, a simple sanity test script is available in `~cs6714/reuters/sanity`. You should run the sanity test script before you submit your solution. To run the sanity test script on a CSE linux machine, simply go inside the folder that contains your `index.py` and `search.py` and type: `~cs6714/reuters/sanity` that will run tests based on examples presented below. Note that it is only a sanity test primarily for formatting, and you are expected to test your project more thoroughly.

The Indexer

Your indexer is run by

```
python3 index.py [folder-of-documents] [folder-of-indexes]
```

where [folder-of-documents] is the path to the directory for the collection of documents to be indexed and [folder-of-indexes] is the path to the directory where the index file(s) should be created. All the files in [folder-of-documents] should be opened as read-only, as you may not have the write permission for these files. If [folder-of-indexes] does not exist, create a new directory as specified. You may create multiple index files although too many index files may slow down your performance. The total size of all your generated index files shall not exceed 20MB (which should be plenty for this project).

The following is an example of how the indexer is run:

```
$ python3 index.py ~cs6714/reuters/data ./MyTestIndex
$
```

The Search

After the indexer is run, your search program is run by

```
python3 search.py [folder-of-indexes]
```

where [folder-of-indexes] is the path to the directory containing the index file(s) that are generated by the indexer. After the above command is executed, it will accept a search query from the standard input and output the result to the standard output as a sequence of document names (the same as their document IDs), one per line and sorted according to their ranking as described in the Ranking section below. It will then continue to accept the search queries from the standard input and output the results to the standard output until the end (i.e., a Ctrl-D). The following example illustrates the required input and output formats:

```
$ python3 search.py ~/Proj/MyTestIndex
Apple
1361
Malaysia
311
908
1356
1675
2956
3051
3438
5169
5195
5216
5258
5285
5382
Australia Technology
3454
271 billions
880
$
```

Ranking

A proximity distance is defined as the number of terms (ignoring decimal numbers when counting) between each pair of terms that match two search terms. The search result is sorted by its minimum sum of the proximity distances between the matching terms (left-to-right by query terms) in an ascending order, then by the number of matching terms occurring in the same order as the search terms, then by the numeric values of the documentIDs (e.g., 72 will be output before 125). For a query with just a single search term, only the third criterion (i.e., ranked by the numeric documentIDs) is employed, since the rest of the criteria are defined based on pairs of search terms. To elaborate this further, consider an example of 4 documents with document IDs 1-4:

<i>DocID</i>	<i>Content</i>
1	apple durian cherry bread egg fennel garlic ham
2	bread garlic ham

```
3      egg bread cherry apple egg fennel ham garlic bread
4      ham garlic bread
```

Given a search query of "garlic bread", all documents contain these two search terms. The minimum proximity distance between the matching terms for these two terms are calculated as below:

Query: garlic bread

Expected output:

DocID	Distance
3	0
4	0
2	0
1	2

For example, in Document 3, *bread* appears twice. We pick the second *bread* to calculate the proximity distance as it will result in a minimum value since it is next to *garlic*, i.e., with a distance 0. We define the matching term that is used to calculate the minimum proximity distance as *the closest matching term*. Although Documents 2, 3, 4 are all having the minimum distance of 0, Documents 3 and 4 rank higher than Document 2 since the two closest matching terms "garlic bread" appear in the same order as the search query, while Document 2 has a different order (i.e., "bread garlic"). Furthermore, since Documents 3 and 4 have the same minimum distance and the closest matching terms are in the same order, they are then sorted by their documentIDs. Finally, Document 1 has the largest minimum distance (2) and hence has the lowest ranking.

Consider another search query "egg ham bread" for the example documents above. Only two documents contain terms matching all three search terms.

Query: egg ham bread

Expected output:

DocID	Distance
3	1+1=2
1	2+3=5

For Document 3, there is more than one term matching *egg* and *bread* respectively. The second *egg* and the second *bread* are chosen as the closest matching terms to achieve the minimum sum of the proximity distances between the 3 matching terms for the 3 search terms. As a result, Document 3 has a smaller minimum distance than Document 1 and hence it ranks higher.

Using the real documents available in `~cs6714/reuters/data`, the following examples present their expected ranked output from your `search.py` and you can study their ranking further by inspecting the content of their corresponding documents.

```
$ python3 search.py ~/Proj/MyTestIndex
australia technology
3454
bank expect distribution
3077
4367
4019
875
$
```

Since the input is from the standard input, your search engine should be able to accept input redirected from a file as shown below:

```
$ cat test1.txt
US finance COMPANY investor
$ python3 search.py ~/Proj/MyTestIndex < test1.txt
5171
3396
3023
5778
1682
$
```

Consider another example of two given documents:

1. butter apple duck chicken
2. chicken butter apple duck

and a search query "apple butter chicken duck".

Both documents have 1 pair in matching order. You may consider it in a similar way to the proximity calculation, i.e. pair-wise, so if apple-butter is in correct order, then +1, etc. For this example:

- butter apple duck chicken: +1 as butter-chicken is correct order
- chicken butter apple duck: +1 as chicken-duck is correct order

Displaying Lines Containing Matching Terms

Given a search query starting with '>' and a space, in addition to the displaying the matching document IDs, the lines of text that contain the closest matching terms are also displayed according to the following rules:

- Each matching document ID is displayed with '>' and a space in front, followed by lines of text that containing the closest matching terms.
- For each matching document, only one line of text that contains its corresponding closest matching term is displayed for each search term.
- Lines of text are displayed according to the order of lines in a document, i.e., output line 1 before line 2.
- In case more than one closest matching term is determined for a search term and they are at different lines, only the first line of these lines is displayed. This includes the case when a search query only consists of one search term such that matching terms can be found in multiple lines (refer to the **apple** example below).

The examples below illustrate some of these rules and the display format.

```
$ python3 search.py ~/Proj/MyTestIndex
> bank expect distribution
> 3077
    The bank said it expects the distribution will be made in
> 4367
    Closing is expected to take place in early April and the
    The partnership will acquire the refining and distribution
    facility with U.S. and foreign banks to finance inventories and
> 4019
    It said it did not know when distributions would be made.
    The bank said it expected to report positive earnings in
> 875
    prepared to negotiate a new distribution based on objective
    bank debt, have increased political pressure on the country to
    The expected drop in prices could result in losses of as
> AUStralia Technology
> 3454
    marketing of high-technology smelting processes invented in
    Australia, notably the Siromelt Zinc Fuming Process.
> apple
> 1361
    The department said stocks of fresh apples in cold storage
$
```

Spelling Correction

You may assume that all search queries shall have matches. For simplicity, we assume there are spelling mistakes in the search terms only when the query does not produce any matches, and your search.py will attempt to correct the spelling with minimum total edit distance to produce matches. During marking, we will only test your program by making spelling mistakes of maximum editing distance of 2 per search term (assuming Insert, Delete and Replace operations and no transpose) and you do not need to correct any mistakes in numeric digits. For your information, search queries that need spelling correction will contribute in total less than 10 points (out of 40 points) of this project.

```
$ python3 search.py ~/Proj/MyTestIndex
australia technologyyy
3454
> auDtralia technologieees
> 3454
marketing of high-technology smelting processes invented in
Australia, notably the Siromelt Zinc Fuming Process.
```

While a generic, lexicon based approach (e.g., based on an English dictionary) shall correct some spelling mistakes, there are cases that it may not work for special terminologies such as:

```
$ python3 search.py ~/Proj/MyTestIndex
> siromell
> 3454
Australia, notably the Siromelt Zinc Fuming Process.
```

Bonus marks: Up to 10% (4 points out of 40 points) bonus marks are allocated for the spelling correction tests when the mis-spelled search terms are actually legit words in the dictionary. For example:

```
$ python3 search.py ~/Proj/MyTestIndex
bank expect distribution
3077
4367
4019
875
bank export distribution
918
875
5116
bank expert distribution
3077
918
4367
4019
875
5116
$
> bark expert distribution
> 3077
    The bank said it expects the distribution will be made in
> 918
    Export-Import Bank of Japan to finance its Pacific Petroleum
    in 1988, is aimed at improving distribution of oil products in
> 4367
    Closing is expected to take place in early April and the
    The partnership will acquire the refining and distribution
    facility with U.S. and foreign banks to finance inventories and
> 4019
    It said it did not know when distributions would be made.
    The bank said it expected to report positive earnings in
> 875
    prepared to negotiate a new distribution based on objective
    "We want to insure that countries receive export quotas
    bank debt, have increased political pressure on the country to
> 1696
    Raffinage-Distribution (CRD) Total France RAFF.PA, told
    journalists that 1986 had marked a return to profit for the
    The subsidiary is expecting to cut its workforce to 6,000
> 5116
    group of North American and Japanese banks to finance the new
    distribution system in 10 U.S. states.
    exports.
```

where expert can be corrected to expect or export. Document 875 matches both of them with different proximity distances. Since we would like to get select the minimum sum, for document 875, "export" instead of "expect" will be chosen for proximity distance calculation to determine the final ranking.

Marking

This project is worth **40 points** (with additional 4 points as bonus, as described above). Your submission will be tested and marked on CSE linux machines using Python3. Therefore, please make sure you have tested your solution on these machines using **Python3** before you submit. You will not receive any marks if your program does not work on CSE linux machines and only works in other environment such as your own laptop. Full marks will be awarded to submissions that follow this specification and pass all the test cases.

Although we do not measure the runtime speed, your indexing program will be terminated if it does not end after **one minute**, and you will receive zero marks for the project (since we cannot get the index generated successfully for further testing); and your search program will be terminated if it does not end after **10 seconds** per search query, and you will receive zero marks for that search query.

Partial Marks

For this project, a search result from your search engine is only considered correct if it contains exactly the same set of document names as the expected answer and their order must be exactly the same as well. We will grant you some partial marks based on F-measure to evaluate how close your result is to the expected answer (in which any correct document names that are in wrong order will be treated as wrong documents). $F\text{-measure} = 2 * (\text{precision} * \text{recall}) / (\text{precision} + \text{recall})$ will be used to calculate the final score for each test when your search results differ from the expected output. The final score for each test will be calculated by: [F-measure round to 2 decimal places] * [fullMarks of that test]. The following examples further illustrate how this scheme works. Given the expected answer Ans, suppose that there are 6 results returned from 6 different search programs.

Ans	Search1	Search2	Search3	Search4	Search5	Search6
1	1	3	2	1	1	1
2	2	2	1	2	2	2
3	3	1	4	3	3	3
4			3	4	5	5
				5	4	6

Their F-measure based on their precision and recall are calculated as below:

```
Search1 - recall: 0.75 precision: 1.00 fmeasure: 0.86
Search2 - recall: 0.25 precision: 0.33 fmeasure: 0.28
Search3 - recall: 0.50 precision: 0.50 fmeasure: 0.50
Search4 - recall: 1.00 precision: 0.80 fmeasure: 0.89
Search5 - recall: 1.00 precision: 0.80 fmeasure: 0.89
Search6 - recall: 0.75 precision: 0.60 fmeasure: 0.67
```

If this search query is worth 2 pts, Search6 will get $0.67 * 2 = 1.34$ pts.

For search queries starting with '>' (the queries with matching lines displayed), no partial marks will be awarded. In order to get full marks, a search result has to exactly match the output of the expected answer (including the document IDs, the matching lines and their order). Therefore, please check your solution with the provided sanity test before submitting the project (and do not leave this till last minute). For your information, search queries starting with '>' will contribute in total less than 10 points (out of 40 points) of this project.

Python3 Libraries

As a requirement of this project, you are not allowed to use any Python libraries other than the Python Standard Library (<https://docs.python.org/3.9/library/index.html>) and NLTK.

For NLTK, you may assume the entire collection (using "all") has been downloaded before the marking process starts. i.e., you should **delete/comment all nltk.download statements**, if any, in your code before your project submission. Note that during your project development, you may want to download only individual required packages rather than "all", as your CSE account has limited space. Please also set the quiet parameter to True in case you forgot to remove the nltk.download statements and the download output affects the auto-marking. For example,

```
nltk.download('package_name', quiet=True)
```

Submission

Deadline: Tuesday 5th November 5:00pm.

The penalty for late submission of assignments will be 5% (of the worth of the assignment) subtracted from the raw mark per day of being late. In other words, earned marks will be lost. **No assignments will be accepted later than 5 days after the deadline.**

Use the give command below to submit the assignment:

```
give cs6714 proj *.py
```

It is **your responsibility** to use classrun to check your submission to make sure that you have submitted all the required files for index.py and search.py to run properly.

```
6714 classrun -check proj
```

Plagiarism

The work you submit must be your own work. Group submissions will not be allowed. Your program must be entirely your own work. Plagiarism detection software will be used to compare all submissions pairwise (including submissions for similar assignments in previous years, if applicable) and serious penalties will be applied, particularly in the case of repeat offences. Your submissions will also be checked against any relevant code available on the Internet. In particular:

- Do not copy ideas or code from others.
- Do not use a publicly accessible repository or allow anyone to see your code.

Please refer to the Academic Honesty and Plagiarism section under Other Useful Information in [the course outline](#) to help you understand what plagiarism is and how it is dealt with at UNSW.

Finally, reproducing, publishing, posting, distributing or translating this assignment is an infringement of copyright and will be referred to UNSW Student Conduct and Integrity for action.

© Copyright 2024, Raymond Wong